

A Survey on Join Algorithms Using MapReduce

Chenhao Fang

University of Wisconsin - Madison

Madison, Wisconsin

cfang45@wisc.edu

ABSTRACT

The popular MapReduce framework is widely used to analyze and process large amount of data. In this work, we present a survey on join algorithms using the MapReduce framework. Join operation is one of the most important and time consuming operations in Database Management Systems. Many research are trying to propose new join algorithms using MapReduce environment. We talk about the MapReduce system and execution process of it. We review the state-of-art algorithms for both two-way and multi-way join algorithms and compare them on several important metrics. We also present analysis of them and discussion about their limitations. This survey can help researchers study the difference of join algorithms in a chronological manner. Concluding, we talk about the direction for future work on join algorithms using MapReduce.

PVLDB Reference Format:

Chenhao Fang. A Survey on Join Algorithms Using MapReduce. PVLDB, 14(1): XXX-XXX, 2020.

doi:XX.XX/XXX.XX

1 INTRODUCTION

The big data and distributing computing is drawing many attentions from both researchers or industries. With the development of new technologies like IoT, AI, cloud computing services and web apps, the amount of data generated every day is at an unprecedented rate. The MapReduce framework was first introduced by Google in [6]. In this programming model and framework for data processing, the details of parallel execution is hidden and users are allowed to only focus on data processing strategies.

Join operation is one of the most important and time consuming operations in Database Management Systems. Therefore, many research are trying to improve the join operation in MapReduce environment using different methods and techniques. In this paper, we will first talk about the MapReduce framework and execution process of MapReduce. Then, we will survey researchers works on join algorithms using MapReduce based on the type of join algorithms using a chronological manner: Map-Side Joins, Reduce-Side Joins, Filter-Based Joins, MapReduce variants and so on. Then, we will analysis and compare them. We summarize the results in a table for quick reference. In the end, we will present our future work about survey and experiments in the future.

This paper is organized as follows: In section 2, we will talk about the background about MapReduce framework and related work. From section 3 to section 7, we will present five categories of join algorithms using MapReduce and make comparison on them. Section 8 and 9 presents the conclusion and proposes future work

2 BACKGROUND

2.1 MapReduce Overview

The MapReduce was first introduced by Jeff Dean and Sanjay Ghemawat in [6]. It is a programming model and the framework implementation that supports the model for processing and generating large data sets. In this programming model, the details of parallel execution is hidden and users are allowed to only focus on data processing strategies. Users only need to specify a *map* function that can process a key-value pair and produce a set of intermediate key/value pairs, and a *reduce* function that merges all the intermediate values associated with the same intermediate key. Users can define their own *map* and *reduce* function based on their own requirements. Some examples that can be implemented easily in MapReduce are: Distributed Grep, Count of URL Access Frequency, Reverse Web-Link Graph, Term-Vector per Host, Inverted Index and Distributed Sort. These examples are discussed in detail in Jeff's original paper.

2.2 Execution

In Jeff and Sanjay's paper [6], the overall flow of a MapReduce operation in Google's implementation is explained in details. They were shown in Figure 1.

If the user calls the MapReduce function, the following procedure happens(The numbers in Figure 1 correspond to the numbers below):

- (1) The user-defined MapReduce program will be split by the MapReduce library into M pieces. The size of each piece can be controlled by user.
- (2) The master node is a special node in MapReduce systems. The master node assigns Map or Reduce work to each worker nodes.
- (3) The workers who got the Map job read the corresponding input split file. It parse the key-value pairs in the input split file and pass them to the user-defined Map function. After this process, the intermediate key-value pairs are stored in buffer.
- (4) The key-value pairs in the buffer are written to disk periodically. They will be partitioned into R parts by partitioning function. The master node will send the locations of these pairs in local disks to the workers who have reduce work.
- (5) The reduce workers will read the buffered data from the local disks. In order to group all the same key together, a

This work is licensed under the Creative Commons BY-NC-ND 4.0 International License. Visit <https://creativecommons.org/licenses/by-nc-nd/4.0/> to view a copy of this license. For any use beyond those covered by this license, obtain permission by emailing info@vldb.org. Copyright is held by the owner/author(s). Publication rights licensed to the VLDB Endowment.

Proceedings of the VLDB Endowment, Vol. 14, No. 1 ISSN 2150-8097.
doi:XX.XX/XXX.XX

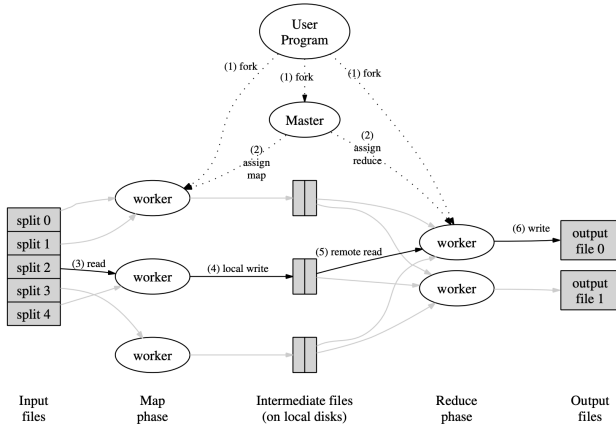


Figure 1: Execution overview of MapReduce. Figure from Jeffery Dean, *MapReduce: Simplified Data Processing on Large Clusters*, 2004.

reduce work need to sort the intermediate key. However, if the size of intermediate data is too large to fit into memory, the reduce worker will use external sort.

- (6) The reduce workers will iterate the sorted data and pass the key and the corresponding set of values to user-defined Reduce function. Then the output of user-defined Reduce function will be appended to a output file for this specific partition.
- (7) When all tasks are completed, the master node will wake up the user program and return back to user code.

Note that the R output files do not need to be combined by the user. The reason is that they will usually be passed to another MapReduce call or use then in some other distributed application which can handle multiple files as input.

2.3 Advantages

The MapReduce is a simple and easy to use model and framework. The major advantages of MapReduce framework for data processing are listed below:

2.3.1 Simple. The MapReduce is extremely easy to use compared with other distributed processing systems. Users only need to define their own Map and Reduce functions and do not need to consider physical distribution of work and execution plans.

2.3.2 Storage Independent. When using MapReduce, you don't need to worry about the underlying storage. The MapReduce is basically independent from storage layers. It can work on different storages like BigTable [5] and so on.

2.3.3 Flexible. The MapReduce framework and model does not dependencies on data model and schema. Therefore, users can easily deal with unstructured data than traditional Database Management Systems.

2.3.4 Scalability. The MapReduce framework has a very high scalability. This is MapReduce's best advantage. 4 years after the

publication time of Jeff's paper [6], Yahoo!'s Hadoop gear could scale out more than 4000 nodes.

2.4 Related Work

In Blanas [4], they provided detailed description and implementation of several join algorithms using the MapReduce framework. For each algorithm, they designed practical pre-processing techniques to improve the performance. Later, they conducted experiments evaluation to compare the different join algorithms on Hadoop cluster. Their work shows the tradeoffs among different join algorithms and the benefits of using MapReduce to process join.

In Lee[7], they provided different technical aspects about the MapReduce framework and classified and summarized works on pros and cons of using MapReduce for data processing. Moreover, the optimization of MapReduce was reviewed. They roughly classified join algorithms into Map-Side Joins and Reduce-Side Joins.

In a paper published in 2020 by Amer[1], they study two-way join algorithms in MapReduce thoroughly. They reviewed state of the art two-way join algorithms in MapReduce. They classified two-way join algorithms into four categories: standard equal-joins, filter-based joins, skew insensitive and MapReduce variants. They summarized their advantages and disadvantages so that the users can choose the most suitable two-way join algorithms in their application.

3 MAP-SIDE JOIN

3.1 Map-Merge Join

The default implementation of map-side join in Hadoop is Map-Merge Join[7]. It is the most common map-side join and works similarly to the sort-merge join in traditional Database Management Systems. There are three phases in Map-Merge Join where the first and second phases are full MapReduce jobs. The Figure 2 shows an example of Map-Merge Join.

In the first phase, the relation R is partitioned and sorted on the join keys, so does relation S in the second phase. In the third phase, mappers read the input partitions and merge them. If the input relations are partitioned and sorted in advance, the third phase can be implemented easily. However, if they are not partitioned and sorted, the overall performance of Map-Merge Join will be degraded.

3.2 Improvements of Map-Merge Join

The Map-Side Partition Merge Join[12] is an improvement of the Map-Merge Join algorithm. The name of this algorithm is quite long and the abbreviation is MSPMJ. It uses on-demand scanning of the second relation and therefore can avoid memory overflow. However, it has the same problem as the Map-Merge Join, which is its ancestor. There are suited for large relation joins, but if at least one of the relations can fit into the node's buffer or cache, they are not optimal.

3.3 Map-Side Replication Join

Given the talk above, the Map-Side Replication Join came out and became the optimal solution [2]. The Map-Side Replication Join

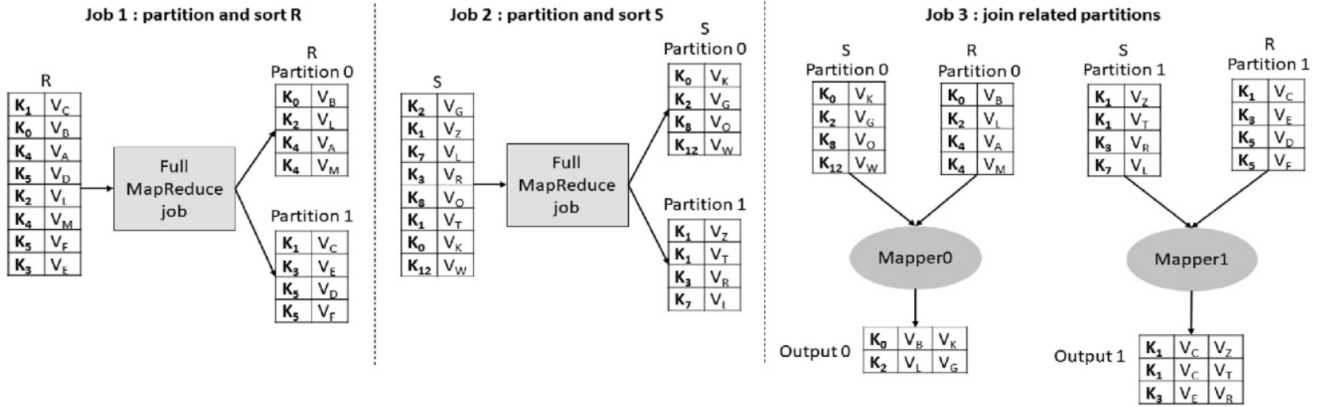


Figure 2: An example of Map-Merge Join. Figure from Amer F. Al-Badarneh, *An analysis of two-way equi-join algorithms under MapReduce*, 2020.

is a map-only join algorithm. The Figure 3 shows the process of Fragment Replicate Join.

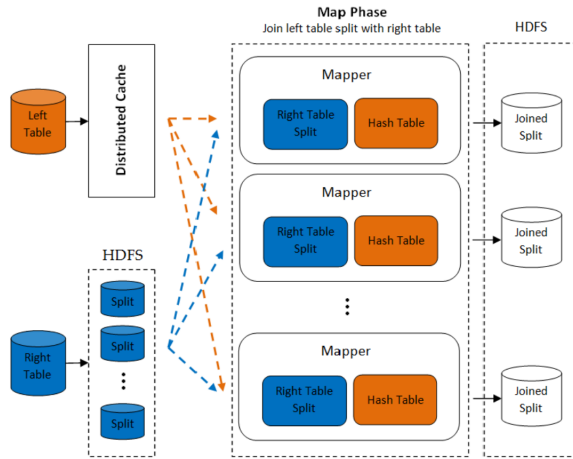


Figure 3: Fragment Replicate Join Implementation. Figure from C. Andreas, *Designing a Parallel Query Engine over Map/Reduce*, 2010.

Consider the left relation L and right relation R, in this case, the size of R is greatly bigger than L. First, the small relation L is replicated into all nodes by using the distributed cache and before the mapper phase, transform it to a hash table. Then for each value in the hash table, the map phase nest an array list to store rows that has the same join key. For each row of the bigger table, search over the unique keys of the small table in hash table. If there's a match, write it to the HDFS output.

However, the Map-Side Partition Merge Join can only be performed if one relation can fit into a node's memory. To improve this, several improved algorithms were proposed, like Reversed Map Join [9] and Broadcast Join [4].

3.4 Improvements of Map-Side Replication Join

The Reversed Map Join were first introduced by Luo and Dong in 2010 [9]. In map phase, the Reversed Map Join builds hash table for each input split of left relation. After that, it iterates through each row of right relation and probes a hash table for each key. Finally, it produces the join result and writes it to the HDFS output.

The Broadcast Join is a combination of the Map-Side Replication Join and the Reversed Map Join. The smaller left relation is broadcast to each mapper in the map phase and kept in each node's memory. It uses hash table to store the smaller relation and find matches by hash table lookup from the other relation. This algorithm avoids I/Os for moving. Also, there's no need to sorting on both relation.

3.5 Comparison

In this subsection, we provide an analysis of different Map-Side Join algorithms based on several different metrics that can affect overall performance in Table 1

4 REDUCE-SIDE JOIN

4.1 Standard Repartition Join

Though Map-Side Join algorithms are easy to implement, the Reduce-Side join are the most common join algorithms when talking about join algorithms in MapReduce. The Standard Repartition Join is the most general Reduce-Side Join algorithms. The Apache Hadoop framework use this algorithm as default implementation. Also, it is widely used in big data systems like Pig, Hive and so on today. This algorithm only run one MapReduce job. One example of Standard Repartition Join is shown in Figure 4.

The special part of this join algorithm is that each mapper will tag each row of relations to identify which relation the row come from. After tagging, the rows that have the same key value are moved to the same reduce node during shuffle and sort phase. And finally, each reduce node will joins the rows which have the same value. This algorithm is similar to the hash join algorithm in traditional Database Management Systems.

Table 1: Comparison of Map-Side Join

Algorithms	Supported Join	Pre-processing	Partitioning	Pros	Cons
Map-Merge Join	Multi-way	Sorting	Hash-based	Easy to implement	Need pre-processing & Sensitive to data skew
MSPMJ	Multi-way	Sorting	Hash-based	Avoid memory overflow	Need pre-processing & Sensitive to data skew
MS Replication Join	Two-way	No	Broadcast	Handle data skew	Has strong condition on relation sizes
Reversed Map Join	Two-way	No	Broadcast	Handle data skew	Complex
Broadcast Join	Two-way	No	Broadcast	Handle data skew	Perform badly if input increases

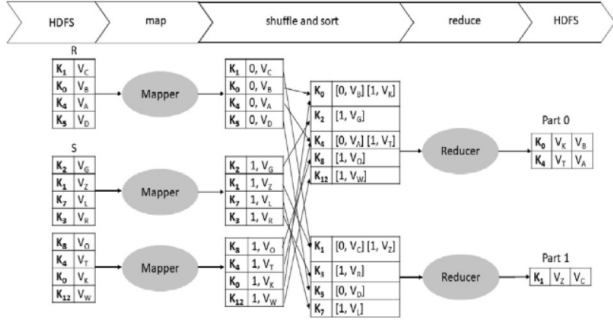


Figure 4: An example of Standard Repartition Join. Figure from Amer F. Al-Badarneh, *An analysis of two-way equi-join algorithms under MapReduce*, 2020.

However, though there's no pre-processing requirements nor special treatments in Standard Repartition Join, it has some drawbacks. It requires extra processing time and may lead to network bottleneck.

4.2 Improved Repartition Join

In order to improve the performance of the Standard Repartition Join, Blanas[4] proposed a new method called Improved Repartition Join. The major different between these two methods are the following three: First, in the map phase, the Improved Repartition Join will be a composition of key and the relation tag. Here, the relation tag will ensure records from Right Relation will sort ahead of records from Left Relation on any given key. The second one is that the hash function is designed in a way that it will only use the join key in the composite key so that the same join key will still be assigned to the same reduce task. Thirdly, due to the reason in the first point, the records that come from Right Relation will be ahead of those from Left Relation. Therefore, only the Right Relation need to be buffered and Left Relation will be streaming data.

4.3 "Schimmy" Reduce-Side Join

Lin proposed a "schimmy" schema to save I/O cost during the process of reduce-side join. [8] In this method, the messages are separated from graph structure data. It will only shuffle the message to save the cost of shuffling data. The mappers only need to emit messages and reducers read graph structure data from HDFS directly. After

these processes, do a reduce-side merge join between the data and messages to perform join.

4.4 Comparison

The comparison of different Reduce-Side Join algorithms based on several different metrics are shown in Table 2

5 FILTER-BASED JOIN

5.1 Semi-Join

In real application, if the Right relation is extremely large, many of its records may not be matched by records from Left relation. A semi-join might be used here to avoid sending unused rows of R over the network.

The semi-join were first implemented by Blanas[4]. This method contains three phases with each phase corresponding to a single MapReduce job. The figure 5 shows an example of semi-join.

The first phase will build a hash table to determine all the unique join keys of Left relation. Then, one reduce node will collect all the unique keys in one file $L.uk$. One assumption is that this file can fit into memory. In the second phase, the $L.uk$ will be loaded into a hash table. The mapper will iterate through Right relation and output the record if the key is in $L.uk$. In the third phase, all the split of R will join with L using the broadcast join method described before.

This semi-join algorithm will reduce the cost of sending unused part of R through network. However, it need to scan an extra round of L. Also, there are still some redundant records in R after the filter that do not join with a particular split of L. This problem can be solved by Per-Split Semi-Join.

5.2 Per-Split Semi-Join

The Per-Split Semi-Join is also proposed by Blanas in the same paper[4]. This method also has three phases like Semi-Join, with each phase corresponding to one MapReduce job. The first phase generates set of unique keys for every split of L_i and store them in several files $L_i.uk$. In the second phase, R will be loaded to hash table in different split. Each $L_i.uk$ file will match the hash table and output R_i . In the third phase, the R_i and L_i will be simply joined by any join method. In the Per-Split Semi-Join, the third phase is extremely easy since the first two phases will do most work. This algorithm can benefit a lot using pre-processing.

However, both Semi-Join and Per-Split Semi-Join will require three phases and three MapReduce jobs. They need to submit jobs many times and therefore will impact the performance.

Table 2: Comparison of Reduce-Side Join

Algorithms	Supported Join	Pre-processing	Partitioning	Pros	Cons
Standard RJ	Multi-way	No	Hash-based	No size restriction	Memory overflow & Sensitive to data skew
Improved RJ	Multi-way	No	No	Fix buffering problem above	Sensitive to data skew
"schimmy" RJ	Two-way	Graph data	No	Save I/O cost	Need complex preprocessing

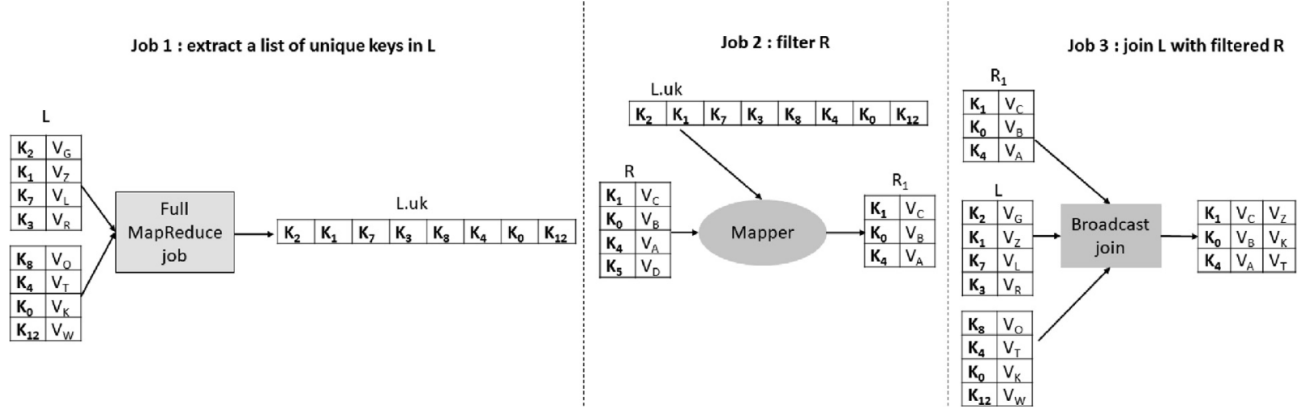


Figure 5: An example of Semi-Join. Figure from Amer F. Al-Badarneh, *An analysis of two-way equi-join algorithms under MapReduce*, 2020.

5.3 Hash-and-Position-Based Per-Split Semi-Join

In 2015, Matono [10] proposed the Hash-and-Position-Based Per-Split Semi-Join to improve the drawback of the above two algorithms. It has a long name, let's just call it HPB PSSJ or HPB. In this method, the author introduced three new reducers and thus number of MapReduce jobs is reduced to two. The first job contains only map phase. The mapper will create hash table $L_i.hp$ for each split of L with the hash value of the keys and their relative positions. In the second phase, the mapper will create hash table H_{Ri} to store the hash values for each split of R and perform actual join.

5.4 Bloom Join

In Zhang [14], the authors proposed a new method of using Bloom filter to join relations in MapReduce. In Bloom Join, the technique is similar to the Semi-Join. They used a Bloom filter based on the R relation to filter part of relation S. In their experiment, the size of both relation R and S is 100 million records.

Two-Stage Join. This algorithm contains two stages and two MapReduce jobs. The Figure 6 shows an example of Two-Stage Bloom Join.

In the first stage, a Bloom filter is build for relation R. Then the mapper uses the Bloom filter to filter relation S and outputs the key-value pairs of the join key a tagged record. All the records will be sent to reducer and the reducer will separate and buffer the input into two sets and then perform a cross-product.

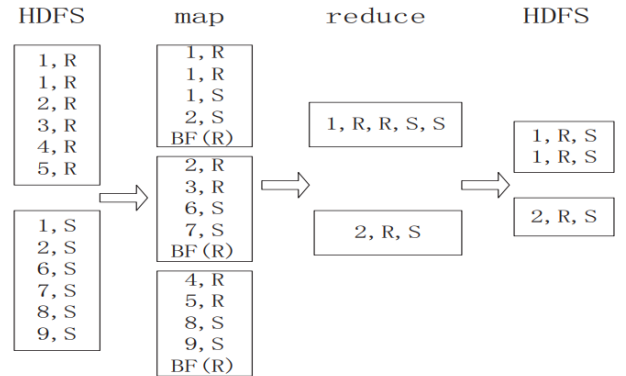


Figure 6: An example of TwoStage Bloom Join Figure from Zhang, *Efficient Processing Distributed Joins with Bloomfilter using MapReduce*, 2013.

Three-Stage Join. This algorithm contains three phases and three MapReduce jobs. The Figure 7 shows an example of Three-Stage Bloom Join.

This first stage is identical to the Two-Stage Join algorithm, while in the second stage, the Bloom filter will be broadcasted to filter S. After this phase, there will be a list of files S_i and they will join with R in the third phase. The jobs in the second and third phase are both map-only jobs.

Table 3: Comparison of Filter-Based-Side Join

Algorithms	Supported Join	Pre-processing	Partitioning	Pros	Cons
Semi-Join	Multi-way	Filtering	Hash-based	Eliminate redundant records	Inefficient if there are too many keys
Per-Split Semi-Join	Multi-way	Filtering	Hash-based	Eliminate redundant records	Complexity
HPB PSSJ	Multi-way	Filtering	Hash-based	Use hash instead of list	More complex
Bloom Join	Two-way	Sorting	No	Space-efficient	Redundant rows in the other table
Two-Stage Join	Two-way	No	Hash-based	Space-efficient	Redundant rows in the other table
Three-Stage Join	Two-way	No	Hash-based	Space-efficient	Need three MapReduce jobs

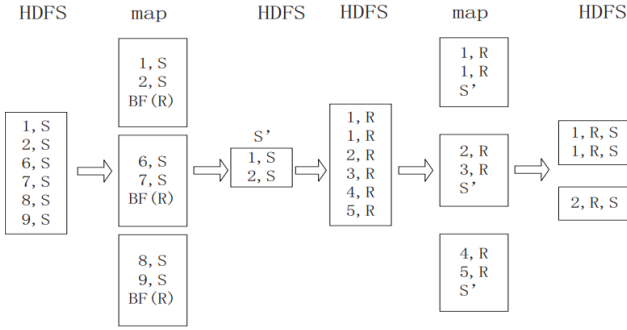


Figure 7: An example of Three-Stage Bloom Join Figure from Zhang, *Efficient Processing Distributed Joins with Bloomfilter using MapReduce*, 2013.

5.5 Comparison

The comparison of different Reduce-Side Join algorithms based on several different metrics are shown in Table 3

6 MAPREDUCE VARIANTS

There have been several research work on modification of the MapReduce framework to cope with the join process.

6.1 Map-Reduce-Merge

Map-Reduce-Merge is the first attempt to optimize join problems in MapReduce framework[13]. To support binary operations including join, the normal Map and Reduce phase in MapReduce framework. In the merge phase, the merger receives the output from two distinguish reduce sources. Then the merge phase processes two key-value pairs and output a new merged key-value pair. The data and control flow of Map-Reduce-Merge framework is shown in Figure 8

This modification retains many advantages of MapReduce, while adding the ability of processing heterogeneous inputs.

6.2 Map-Join-Reduce

Map-Join-Reduce is another variant of MapReduce framework that can perform one-phase join. It adds a join stage before reduce stage. First, every mapper reads records from relations that will be joined. Then, the mapper outputs are shuffled and moved to joiners. The joiners are run inside the reduce nodes. The joiners will perform actual join operations.

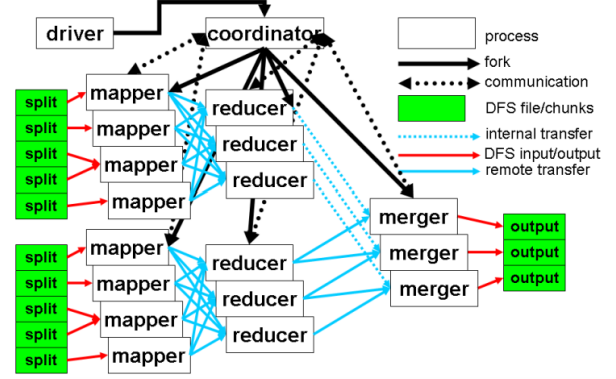


Figure 8: Data and control flow for the MapReduce-Merge framework. Figure from Yang, *Map-Reduce-Merge: Simplified Relational Data Processing on Large Clusters*, 2013.

6.3 Comparison

Though MapReduce variants can help the framework improve the ability of data processing, they still require a lot of spaces to store the intermediate results. Also, they need to do modification to a mature framework, therefore, they were not used widely compared with algorithms talked above.

The comparison of different MapReduce variants are shown in Table 4

7 OTHER JOIN ALGORITHMS

7.1 Hybrid Hadoop Join

Atta[3] developed a hybrid approach that is a combination of Map-Side and Reduce-Side join. This algorithm only performs a partial partition-wise join and the join happens in reduce phase. The ?? shows the dataflow of Hybrid Hadoop Join.

The Hybrid Hadoop join has two phases and two MapReduce jobs. In the first job, the algorithm partitions and sorts the smaller relation to different splits R_i . Next, in the second job, each mapper will iterate through rows of R_i and pass the key-value pairs to the partitioning function. It will partition and sort S in the same way that R was partitioned and sorted. Then the output of partitioning function will be passed to reducers and in each reducer, the counter part of R_i will be loaded into a hash table. Finally, the reducer will

Table 4: Comparison of MapReduce variants

Algorithms	Supported Join	Pre-processing	Partitioning	Pros	Cons
Map-Reduce-Merge	Two-way	No	No	Efficient for heterogenous data	Hard to implement
Map-Join-Reduce	Two-way	No	No	Efficient for heterogenous data	Hard to implement

Table 5: Comparison of MapReduce variants

Algorithms	Supported Join	Pre-processing	Partitioning	Pros	Cons
Hybrid Hadoop Join	Multi-way	sorting	Hash-based	Less prerequisites	Complex
Theta Join	Multi-way	No	No	Minimize running time	Need partitioning and replication

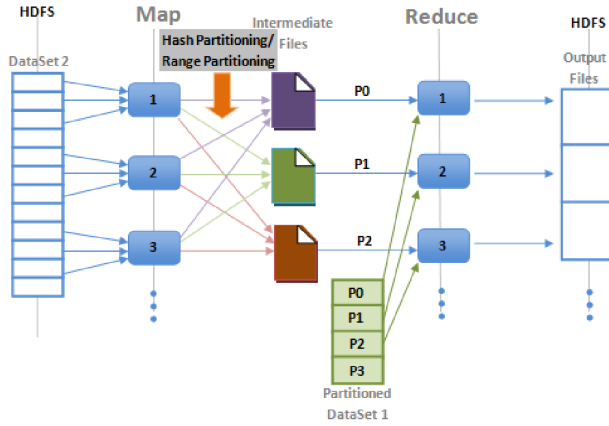


Figure 9: Data and control flow for the Hybrid Hadoop framework. Figure from Atta, *Implementation and Analysis of Join Algorithms to Handle Skew for the Hadoop Map/Reduce Framework*, 2010.

probe the hash table for each input key of S_i and perform the join operation. The final output will be written to HDFS.

This algorithm has the advantage of Map-Side Join and Reduce-Side Join. Moreover, it eliminates the overhead of the tagging, and since only one relation needs to be partitioned, it limits the pre-processing. However, it still has drawbacks of the Map-Side Join and Reduce-Side Join. Also, the parameters have to be set properly in this algorithm.

7.2 Theta Join

The first Theta Join using MapReduce is first introduced by Okcan[11]. The main idea of this algorithm is to reduce task by partitioning join space, through balancing the input and output cost of reduce jobs. Also, it will minimize the number of input records that are duplicated. This algorithm uses a reducer-centered cost model that can calculate the total cost of cross product of mapper output. Using this model, it just assigns the mapper output to reducers that can minimize the job running time.

7.3 Comparison

The comparison between these two algorithms is shown in Table 5.

8 CONCLUSION

The join operation is the most important and costly operation in Database Management Systems. In this survey, we reviewed several state-of-the-art join algorithms in the MapReduce framework and compared them using several important metrics. We first discussed the MapReduce system and its methodology. The core of this work is surveying on five different kinds of join algorithms: Map-Side Join, Reduce-Side Join, Filter-Based Join, MapReduce Variants and other types. We hope this work will help researchers compare different join algorithms using MapReduce easily when they are trying to propose new improvements or new ideas in this area.

In the future work, we will research on other join algorithms and different join conditions like Similarity Join, KNN Join, Top-K Join and so on. We also plan to run experiments on the Hadoop platform or modern Cloud Computing platforms to test the efficiency of these algorithms. We will provide a cost-based comparison after the experiments. Evaluating the join processing in SparkSQL is also an interesting topic for further investigation.

9 APPENDIX

The comparison of all the Join algorithms using MapReduce based on several different metrics were shown in 6

REFERENCES

- [1] Amer F Al-Badarnah and Salahaldeen Atef Rababa. 2020. An Analysis of Two-Way Equi-Join Algorithms Under MapReduce. *Journal of King Saud University-Computer and Information Sciences* (2020).
- [2] Chatzistergiou Andreas. 2010. *Designing a Parallel Query Engine over Map/Reduce*. Ph.D. Dissertation. Citeseer.
- [3] Fariha Atta. 2010. *Implementation and analysis of join algorithms to handle skew for the Hadoop Map/Reduce framework*. Ph.D. Dissertation. Master's thesis, MSc Informatics, School of Informatics, University of Edinburgh.
- [4] Spyros Blanas, Jignesh M Patel, Vuk Ercegovic, Jun Rao, Eugene J Shekita, and Yuanyuan Tian. 2010. A comparison of join algorithms for log processing in mapreduce. In *Proceedings of the 2010 ACM SIGMOD International Conference on Management of data*. 975–986.
- [5] Fay Chang, Jeffrey Dean, Sanjay Ghemawat, Wilson C. Hsieh, Deborah A. Wal-lach, Mike Burrows, Tushar Chandra, Andrew Fikes, and Robert E. Gruber. 2006. Bigtable: A Distributed Storage System for Structured Data. In *7th USENIX Symposium on Operating Systems Design and Implementation (OSDI)*. 205–218.
- [6] Jeffrey Dean and Sanjay Ghemawat. 2004. MapReduce: Simplified Data Processing on Large Clusters. In *OSDI'04: Sixth Symposium on Operating System Design and Implementation*. San Francisco, CA, 137–150.

- [7] Kyong-Ha Lee, Yoon-Joon Lee, Hyunsik Choi, Yon Dohn Chung, and Bongki Moon. 2012. Parallel data processing with MapReduce: a survey. *AcM SIGMOD Record* 40, 4 (2012), 11–20.
- [8] Jimmy Lin and Chris Dyer. 2010. Data-intensive text processing with MapReduce. *Synthesis Lectures on Human Language Technologies* 3, 1 (2010), 1–177.
- [9] Gang Luo and Liang Dong. 2010. Adaptive join plan generation in hadoop. *Duke University, Durham NC, USA* (2010).
- [10] Akiyoshi Matono, Hirotaka Ogawa, and Isao Kojima. 2015. Improvement of Join Algorithms for Low-Selectivity Joins on MapReduce. In *Australasian Database Conference*. Springer, 117–128.
- [11] Alper Okcan and Mirek Riedewald. 2011. Processing theta-joins using MapReduce. In *Proceedings of the 2011 ACM SIGMOD International Conference on Management of data*. 949–960.
- [12] A Pigul. 2012. Comparative Study parallel join algorithms for MapReduce environment. 23 (2012).
- [13] Hung-chih Yang, Ali Dasdan, Ruey-Lung Hsiao, and D. Stott Parker. 2007. MapReduce-Merge: Simplified Relational Data Processing on Large Clusters. In *Proceedings of the 2007 ACM SIGMOD International Conference on Management of Data* (Beijing, China) (*SIGMOD '07*). Association for Computing Machinery, New York, NY, USA, 1029–1040. <https://doi.org/10.1145/1247480.1247602>
- [14] Changchun Zhang, Lei Wu, and Jing Li. 2013. Efficient processing distributed joins with bloomfilter using mapreduce. *International Journal of Grid and Distributed Computing* 6, 3 (2013), 43–58.

Table 6: Comparison of Join Algorithms using MapReduce

Algorithms	Supported Join	Pre-processing	Partitioning	Pros	Cons
Map-Merge Join	Multi-way	Sorting	Hash-based	Easy to implement	Sensitive to data skew
MSPMJ	Multi-way	Sorting	Hash-based	Avoid memory overflow	Sensitive to data skew
MS Replication Join	Two-way	No	Broadcast	Handle data skew	Has strong condition on relation sizes
Reversed Map Join	Two-way	No	Broadcast	Handle data skew	Complex
Broadcast Join	Two-way	No	Broadcast	Handle data skew	Perform badly if input increases
Standard RJ	Multi-way	No	Hash-based	No size restriction	Memory overflow
Improved RJ	Multi-way	No	No	Fix buffering problem above	Sensitive to data skew
"schimmy" RJ	Two-way	Graph data	No	Save I/O cost	Need complex preprocessing
Semi-Join	Multi-way	Filtering	Hash-based	Eliminate redundant records	Inefficient if there are too many keys
Per-Split Semi-Join	Multi-way	Filtering	Hash-based	Eliminate redundant records	Complexity
HPB PSSJ	Multi-way	Filtering	Hash-based	Use hash instead of list	More complex
Bloom Join	Two-way	Sorting	No	Space-efficient	Redundant rows in the other table
Two-Stage Join	Two-way	No	Hash-based	Space-efficient	Redundant rows in the other table
Three-Stage Join	Two-way	No	Hash-based	Space-efficient	Need three MapReduce jobs
Map-Reduce-Merge	Two-way	No	No	Efficient for heterogenous data	Hard to implement
Map-Join-Reduce	Two-way	No	No	Efficient for heterogenous data	Hard to implement
Hybrid Hadoop Join	Multi-way	sorting	Hash-based	Less prerequisites	Complex
Theta Join	Multi-way	No	No	Minimize running time	Need partitioning and replication